

Laboratory Practice III – Practical 1

Name: Anuj Sachin Dhole
Roll No: B21042
Class: BE CEA
Subject: Laboratory Practice III (Machine Learning)

Practical 1

Problem Statement:

Predict the price of the Uber ride from a given pickup point to the agreed drop-off location.

Tasks to Perform:

- Pre-process the dataset.
- Identify outliers.
- Check the correlation.
- Implement Linear Regression and Random Forest Regression models.
- Evaluate the models and compare their respective scores like R², RMSE, etc.

Dataset:

Source: [Uber Fares Dataset on Kaggle](#)

```
In [126]: # Predict the price of the Uber ride from a given pickup point to the agreed drop-off location.  
# Perform following tasks:  
# 1. Pre-process the dataset.  
# 2. Identify outliers.  
# 3. Check the correlation.  
# 4. Implement linear regression and random forest regression models.  
# 5. Evaluate the models and compare their respective scores like R2, RMSE, etc.  
# Dataset link: https://www.kaggle.com/datasets/yasserh/uber-fares-dataset
```

```
In [1]: # 1. Load libraries and dataset  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression  
from sklearn.ensemble import RandomForestRegressor  
from sklearn.metrics import r2_score, mean_squared_error
```

```
In [2]: import os  
print("Current directory:", os.getcwd())
```

```
# List files in current directory  
print("*files here*", os.listdir())
```

```
Current directory: C:\Users\MBAL PC-17\Desktop\Anuj Dhole
```

```
Files here: ['ipython_checkpoints', 'Prac1.ipynb', 'uber.csv']
```

```
In [3]: df = pd.read_csv('uber.csv')  
df.dropna(inplace=True)
```

```
In [4]: print("Data sample:")  
print(df.head())
```

```
print("\nData info:")  
print(df.info())
```

```
print("\nMissing values:")  
print(df.isnull().sum())
```

```
Data sample:  
Unnamed: 0 key fare_amount \
```

```
0 24238194 2015-05-07 19:52:06.0000003 7.5
```

```
1 27835199 2009-07-17 20:04:56.0000002 7.7
```

```
2 44984355 2009-08-24 21:45:00.00000061 12.9
```

```
3 25894730 2009-06-26 08:22:21.0000001 5.3
```

```
4 17610152 2014-08-28 17:47:00.000000188 16.0
```

```
dropoff_longitude dropoff_latitude passenger_count \
```

```
0 -73.999512 40.723217 1
```

```
1 -73.994710 40.750325 1
```

```
2 -73.962565 40.77647 1
```

```
3 -73.965316 40.803349 3
```

```
4 -73.973082 40.761247 5
```

```
Data info:  
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 199999 entries, 0 to 199999
```

```
Data columns (total 9 columns):
```

```
Unnamed: 0 199999 non-null int64
```

```
key 199999 non-null object
```

```
fare_amount 199999 non-null float64
```

```
pickup_datetime 199999 non-null object
```

```
pickup_longitude 199999 non-null float64
```

```
pickup_latitude 199999 non-null float64
```

```
dropoff_longitude 199999 non-null float64
```

```
dropoff_latitude 199999 non-null float64
```

```
passenger_count 199999 non-null int64
```

```
dtypes: float64(5), int64(2), object(2)
```

```
memory usage: 15.3+ MB
```

```
None
```

```
Missing values:
```

```
Unnamed: 0 0
```

```
key 0
```

```
fare_amount 0
```

```
pickup_datetime 0
```

```
pickup_longitude 0
```

```
pickup_latitude 0
```

```
dropoff_longitude 0
```

```
dropoff_latitude 0
```

```
passenger_count 0
```

```
dtype: int64
```

```
In [5]: # Convert pickup_datetime to datetime format  
df['pickup_datetime'] = pd.to_datetime(df['pickup_datetime'], errors='coerce')  
# converts the pickup_datetime column from strings into actual datetime objects
```

```
# pd.to_datetime() tries to parse each value in df['pickup_datetime'] and turn it into a datetime64 type (date + time).
```

```
# errors='coerce' means:
```

```
# If a value can't be parsed into a valid date/time, instead of raising an error, it will convert that value into a missing value (NaN).
```

```
In [6]: df.head()
```

```
Out[6]:
```

Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count
0	24238194	19:52:06.0000003	2015-05-07 19:52:06.000000	-73.999817	40.738354	-73.999512	40.723217	1

```
0 24238194 19:52:06.0000003 2015-05-07 19:52:06.000000 -73.999817 40.738354 -73.999512 40.723217 1
```

```
1 27835199 20:04:56.0000002 2009-07-17 20:04:56.000000 -73.994355 40.728225 -73.994710 40.750325 1
```

```
2 44984355 21:45:00.00000061 2009-08-24 21:45:00.000000 -74.005043 40.740770 -73.962565 40.772647 1
```

```
3 25894730 08:22:21.0000001 2009-06-26 08:22:21.000000 -73.976124 40.790844 -73.965316 40.803349 3
```

```
4 17610152 17:47:00.000000188 2014-08-28 17:47:00.000000 -73.925023 40.744085 -73.973082 40.761247 5
```

```
In [7]: # Extract datetime features  
df['hour'] = df['pickup_datetime'].dt.hour  
df['day'] = df['pickup_datetime'].dt.day  
df['weekday'] = df['pickup_datetime'].dt.weekday  
df['month'] = df['pickup_datetime'].dt.month
```

```
In [8]: df.head()
```

```
Out[8]:
```

Unnamed: 0	key	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	passenger_count	hour
0	24238194	19:52:06.0000003	2015-05-07 19:52:06.000000	-73.999817	40.738354	-73.999512	40.723217	1	19

```
0 24238194 19:52:06.0000003 2015-05-07 19:52:06.000000 -73.999817 40.738354 -73.999512 40.723217 1
```

```
1 27835199 20:04:56.0000002 2009-07-17 20:04:56.000000 -73.994355 40.728225 -73.994710 40.750325 1
```

```
2 44984355 21:45:00.00000061 2009-08-24 21:45:00.000000 -74.005043 40.740770 -73.962565 40.772647 1
```

```
3 25894730 08:22:21.0000001 2009-06-26 08:22:21.000000 -73.976124 40.790844 -73.965316 40.803349 3
```

```
4 17610152 17:47:00.000000188 2014-08-28 17:47:00.000000 -73.925023 40.744085 -73.973082 40.761247 5
```

```
In [9]: # Haversine distance function  
def haversine_distance(lat1, lon1, lat2, lon2):  
    R = 6371 # Earth's radius in km  
    phi1 = np.radians(lat1)  
    phi2 = np.radians(lat2)  
    d_phi = np.radians(lon2 - lon1)  
    d_lambda = np.radians(lon2 - lon1)  
  
    a = np.sin(d_phi / 2)**2 +  
        np.cos(phi1) * np.cos(phi2) * np.sin(d_lambda / 2)**2  
  
    c = 2 * np.arcsin(np.sqrt(a))  
    return R * c
```

```
# Apply Haversine formula to get distance
```

```
df['trip_distance_km'] = haversine_distance(df['pickup_latitude'], df['pickup_longitude'], df['dropoff_latitude'], df['dropoff_longitude'])
```

```
df['trip_distance_km'] = df['trip_distance_km']**2
```

```
# Alternative to Haversine  
# Didn't use Haversine as it is complicated
```

```
# This Works well if your coordinates are in a flat (planar) coordinate system.
```

```
# Haversine Formula calculates the great-circle distance between two points on the Earth's surface (accounts for Earth's curvature).
```

```
In [11]: # Filter invalid values (NOT Outlier removal, just sanity check)  
df = df[(df['fare_amount'] > 0) & (df['fare_amount'] < 100)]  
df = df[(df['passenger_count'] > 0) & (df['passenger_count'] < 6)]  
df = df[(df['trip_distance_km'] > 0)]
```

```
In [12]: # 2. Outlier detection using IQR  
def detect_outliers_iqr(data, column):  
    Q1 = data[column].quantile(0.25)  
    Q3 = data[column].quantile(0.75)  
    IQR = Q3 - Q1  
    lower_bound = Q1 - 1.5 * IQR  
    upper_bound = Q3 + 1.5 * IQR  
    outliers = data[(data[column] < lower_bound) | (data[column] > upper_bound)]  
  
    print(f"\nOutliers in '{column}'")  
    print(f"Q1 = {Q1}, Q3 = {Q3}, IQR = {IQR:.2f}")  
    print(f"Lower Bound = {lower_bound:.2f}, Upper Bound = {upper_bound:.2f}")  
    print(f"Number of Outliers: {len(outliers)}")
```

```
# Boxplot
```

```
plt.figure(figsize=(6, 4))  
sns.boxplot(x=column)
```

```
plt.title(f'Boxplot for {column}')  
plt.show()
```

```
return outliers
```

```
outlier_columns = ['fare_amount', 'trip_distance_km', 'passenger_count']
```

```
for col in outlier_columns:  
    detect_outliers_iqr(df, col)
```

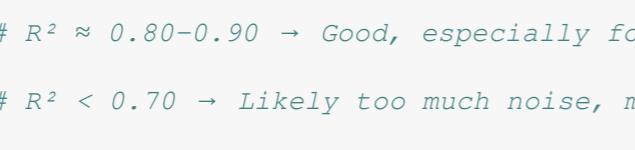
```
Outliers in 'fare_amount':
```

```
Q1 = 1.28, Q3 = 3.96, IQR = 2.68
```

```
Lower Bound = -2.73, Upper Bound = 7.97
```

```
Number of Outliers: 16544
```

```
Boxplot for fare_amount
```



```
Outliers in 'trip_distance_km':
```

```
Q1 = 1.00, Q3 = 2.00, IQR = 1.00
```

```
Lower Bound = -0.50, Upper Bound = 3.50
```

```
Number of Outliers: 21908
```

```
Boxplot for trip_distance_km
```



```
Outliers in 'passenger_count':
```

```
Q1 = 1.00, Q3 = 2.00, IQR = 1.00
```

```
Lower Bound = -0.50, Upper Bound = 3.50
```

```
Number of Outliers: 16544
```

```
Boxplot for passenger_count
```



```
In [13]: def remove_outliers_iqr(data, column):  
    Q1 = data[column].quantile(0.25)  
    Q3 = data[column].quantile(0.75)  
    IQR = Q3 - Q1  
    lower_bound = Q1 - 1.5 * IQR  
    upper_bound = Q3 + 1.5 * IQR  
    outliers = data[(data[column] < lower_bound) | (data[column] > upper_bound)]  
  
    print(f"\nOutliers in '{column}'")  
    print(f"Q1 = {Q1}, Q3 = {Q3}, IQR = {IQR:.2f}")  
    print(f"Lower Bound = {lower_bound:.2f}, Upper Bound = {upper_bound:.2f}")  
    print(f"Number of Outliers: {len(outliers)}")
```

```
# Boxplot
```

```
plt.figure(figsize=(6, 4))  
sns.boxplot(x=column)
```

```
plt.title(f'Boxplot for {column}')  
plt.show
```